

Introduction to numpy:

Package for scientific computing with Python

Numerical Python, or "Numpy" for short, is a foundational package on which many of the most common data science packages are built. Numpy provides us with high performance multi-dimensional arrays which we can use as vectors or matrices.

The key features of numpy are:

- **ndarrays:** n-dimensional arrays of the same data type which are fast and space-efficient. There are a number of built-in methods for ndarrays which allow for rapid processing of data without using loops (e.g., compute the mean).
- **Broadcasting:** a useful tool which defines implicit behavior between multi-dimensional arrays of different sizes.
- **Vectorization:** enables numeric operations on ndarrays.
- **Input/Output:** simplifies reading and writing of data from/to file.

Additional Recommended Resources:

[Numpy Documentation \(https://docs.scipy.org/doc/numpy/reference/\)](https://docs.scipy.org/doc/numpy/reference/)

Python for Data Analysis by Wes McKinney

Python Data science Handbook by Jake VanderPlas

Getting started with ndarray

ndarrays are time and space-efficient multidimensional arrays at the core of numpy. Like the data structures in Week 2, let's get started by creating ndarrays using the numpy package.

How to create Rank 1 numpy arrays:

```
In [3]: import numpy as np

an_array = np.array([3, 33, 333]) # Create a rank 1 array

print(type(an_array))           # The type of an ndarray is: "<class 'numpy.ndarray'>"

<class 'numpy.ndarray'>
```

```
In [4]: # test the shape of the array we just created, it should have just one dimension (Rank 1)
print(an_array.shape)

(3,)
```

```
In [5]: # because this is a 1-rank array, we need only one index to access each element
print(an_array[0], an_array[1], an_array[2])

3 33 333
```

```
In [6]: an_array[0] = 888          # ndarrays are mutable, here we change an element of the array

print(an_array)

[888  33 333]
```

How to create a Rank 2 numpy array:

A rank 2 **ndarray** is one with two dimensions. Notice the format below of `[[row] , [row]]`. 2 dimensional arrays are great for representing matrices which are often useful in data science.

```
In [7]: another = np.array([[11,12,13],[21,22,23]]) # Create a rank 2 array

print(another) # print the array

print("The shape is 2 rows, 3 columns: ", another.shape) # rows x columns

print("Accessing elements [0,0], [0,1], and [1,0] of the ndarray: ",
      another[0, 0], ", ", another[0, 1], ", ", another[1, 0])

[[11 12 13]
 [21 22 23]]
The shape is 2 rows, 3 columns: (2, 3)
Accessing elements [0,0], [0,1], and [1,0] of the ndarray: 11 , 12 , 21
```

There are many way to create numpy arrays:

Here we create a number of different size arrays with different shapes and different pre-filled values. numpy has a number of built in methods which help us quickly and easily create multidimensional arrays.

```
In [31]: import numpy as np

# create a 2x2 array of zeros
ex1 = np.zeros((3,3))
print(ex1)
print("-----")
print(ex1[:2,])
```

```
[[ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]]
-----
[[ 0.  0.  0.]
 [ 0.  0.  0.]]
```

```
In [9]: # create a 2x2 array filled with 9.0
ex2 = np.full((2,2), 9.0)
print(ex2)
```

```
[[ 9.  9.]
 [ 9.  9.]]
```

```
In [10]: # create a 2x2 matrix with the diagonal 1s and the others 0
ex3 = np.eye(2,2)
print(ex3)
```

```
[[ 1.  0.]
 [ 0.  1.]]
```

```
In [11]: # create an array of ones
ex4 = np.ones((1,2))
print(ex4)
```

```
[[ 1.  1.]]
```

```
In [12]: # notice that the above ndarray (ex4) is actually rank 2, it is a 2x1 array
print(ex4.shape)
```

```
# which means we need to use two indexes to access an element
print()
print(ex4[0,1])
```

```
(1, 2)
```

```
1.0
```

```
In [ ]: # create an array of random floats between 0 and 1
ex5 = np.random.random((2,2))
print(ex5)
```

Array Indexing

Slice indexing:

Similar to the use of slice indexing with lists and strings, we can use slice indexing to pull out sub-regions of ndarrays.

```
In [13]: import numpy as np

# Rank 2 array of shape (3, 4)
an_array = np.array([[11,12,13,14], [21,22,23,24], [31,32,33,34]])
print(an_array)

[[11 12 13 14]
 [21 22 23 24]
 [31 32 33 34]]
```

Use array slicing to get a subarray consisting of the first 2 rows x 2 columns.

```
In [19]: a_slice = np.array(an_array[:2, 1:3])
print(a_slice)

[[1000  13]
 [ 22  23]]
```

When you modify a slice, you actually modify the underlying array.

```
In [21]: print("Before:", an_array[0, 1])  #inspect the element at 0, 1
a_slice[0, 0] = 12  # a_slice[0, 0] is the same piece of data as an_array[0, 1]
print("After:", an_array[0, 1])

Before: 1000
After: 1000
```

Use both integer indexing & slice indexing

We can use combinations of integer indexing and slice indexing to create different shaped matrices.

```
In [22]: # Create a Rank 2 array of shape (3, 4)
an_array = np.array([[11,12,13,14], [21,22,23,24], [31,32,33,34]])
print(an_array)

[[11 12 13 14]
 [21 22 23 24]
 [31 32 33 34]]
```

```
In [23]: # Using both integer indexing & slicing generates an array of lower rank
row_rank1 = an_array[1, :] # Rank 1 view

print(row_rank1, row_rank1.shape) # notice only a single []

[21 22 23 24] (4,)
```

```
In [24]: # Slicing alone: generates an array of the same rank as the an_array
row_rank2 = an_array[1:2, :] # Rank 2 view

print(row_rank2, row_rank2.shape) # Notice the [[ ]]

[[21 22 23 24]] (1, 4)
```

```
In [25]: #We can do the same thing for columns of an array:

print()
col_rank1 = an_array[:, 1]
col_rank2 = an_array[:, 1:2]

print(col_rank1, col_rank1.shape) # Rank 1
print()
print(col_rank2, col_rank2.shape) # Rank 2

[12 22 32] (3,)

[[12]
 [22]
 [32]] (3, 1)
```

Array Indexing for changing elements:

Sometimes it's useful to use an array of indexes to access or change elements.

```
In [26]: # Create a new array
an_array = np.array([[11,12,13], [21,22,23], [31,32,33], [41,42,43]])

print('Original Array:')
print(an_array)
```

```
Original Array:
[[11 12 13]
 [21 22 23]
 [31 32 33]
 [41 42 43]]
```

```
In [27]: # Create an array of indices
col_indices = np.array([0, 1, 2, 0])
print('\nCol indices picked : ', col_indices)

row_indices = np.arange(4)
print('\nRows indices picked : ', row_indices)
```

```
Col indices picked : [0 1 2 0]
```

```
Rows indices picked : [0 1 2 3]
```

```
In [28]: # Examine the pairings of row_indices and col_indices.  These are the elements
          we'll change next.
for row,col in zip(row_indices,col_indices):
    print(row, ", ", col)
```

```
0 , 0
1 , 1
2 , 2
3 , 0
```

```
In [29]: # Select one element from each row
print('Values in the array at those indices: ',an_array[row_indices, col_indices])
```

```
Values in the array at those indices: [11 22 33 41]
```

```
In [ ]: # Change one element from each row using the indices selected
an_array[row_indices, col_indices] += 100000

print('\nChanged Array:')
print(an_array)
```

Boolean Indexing

Array Indexing for changing elements:

```
In [ ]: # create a 3x2 array
an_array = np.array([[11,12], [21, 22], [31, 32]])
print(an_array)
```

```
In [ ]: # create a filter which will be boolean values for whether each element meets
        this condition
filter = (an_array > 15)
filter
```

Notice that the filter is a same size ndarray as an_array which is filled with True for each element whose corresponding element in an_array which is greater than 15 and False for those elements whose value is less than 15.

```
In [ ]: # we can now select just those elements which meet that criteria
print(an_array[filter])
```

```
In [ ]: # For short, we could have just used the approach below without the need for t
        he separate filter array.

an_array[(an_array % 2 == 0)]
```

What is particularly useful is that we can actually change elements in the array applying a similar logical filter. Let's add 100 to all the even values.

```
In [ ]: an_array[an_array % 2 == 0] +=100
print(an_array)
```

Datatypes and Array Operations

Datatypes:

```
In [ ]: ex1 = np.array([11, 12]) # Python assigns the data type
print(ex1.dtype)
```

```
In [ ]: ex2 = np.array([11.0, 12.0]) # Python assigns the data type
print(ex2.dtype)
```

```
In [ ]: ex3 = np.array([11, 21], dtype=np.int64) #You can also tell Python the data t
        ype
print(ex3.dtype)
```

```
In [ ]: # you can use this to force floats into integers (using floor function)
ex4 = np.array([11.1,12.7], dtype=np.int64)
print(ex4.dtype)
print()
print(ex4)
```

```
In [ ]: # you can use this to force integers into floats if you anticipate
# the values may change to floats later
ex5 = np.array([11, 21], dtype=np.float64)
print(ex5.dtype)
print()
print(ex5)
```

Arithmetic Array Operations:

```
In [ ]: x = np.array([[111,112],[121,122]], dtype=np.int)
y = np.array([[211.1,212.1],[221.1,222.1]], dtype=np.float64)

print(x)
print()
print(y)
```

```
In [ ]: # add
print(x + y)          # The plus sign works
print()
print(np.add(x, y))   # so does the numpy function "add"
```

```
In [ ]: # subtract
print(x - y)
print()
print(np.subtract(x, y))
```

```
In [ ]: # multiply
print(x * y)
print()
print(np.multiply(x, y))
```

```
In [ ]: # divide
print(x / y)
print()
print(np.divide(x, y))
```

```
In [ ]: # square root
print(np.sqrt(x))
```

```
In [ ]: # exponent (e ** x)
print(np.exp(x))
```


Statistical Methods, Sorting, and Set Operations:

Basic Statistical Operations:

```
In [33]: # setup a random 2 x 4 matrix
arr = 10 * np.random.randn(2,5)
print(arr)

[[-13.29415778  6.58124512 -18.04621171  5.76197528 13.49582377]
 [-29.93674421  2.30901788  0.73468679 -19.10746865 -5.95183744]]
```

```
In [34]: # compute the mean for all elements
print(arr.mean())

-5.74536709543
```

```
In [35]: # compute the means by row
print(arr.mean(axis = 1))

[ -1.10026506 -10.39046913]
```

```
In [36]: # compute the means by column
print(arr.mean(axis = 0))

[-21.615451  4.4451315 -8.65576246 -6.67274668  3.77199316]
```

```
In [37]: # sum all the elements
print(arr.sum())

-57.4536709543
```

```
In [39]: # compute the medians
print(np.median(arr,axis=0))
#important- np.ndarray has no attribute median - call directly from the np
library and pass the array as the argument

[-21.615451  4.4451315 -8.65576246 -6.67274668  3.77199316]
```

Sorting:

```
In [40]: # create a 10 element array of randoms
         unsorted = np.random.randn(10)

         print(unsorted)

[ 0.06725739  0.32812902  0.79000774  1.43855797 -0.00846992 -1.25954441
  1.01884829  0.62523893  2.59885886  0.23700386]
```

```
In [41]: # create copy and sort
         sorted = np.array(unsorted)
         sorted.sort()

         print(sorted)
         print()
         print(unsorted)

[-1.25954441 -0.00846992  0.06725739  0.23700386  0.32812902  0.62523893
  0.79000774  1.01884829  1.43855797  2.59885886]

[ 0.06725739  0.32812902  0.79000774  1.43855797 -0.00846992 -1.25954441
  1.01884829  0.62523893  2.59885886  0.23700386]
```

```
In [42]: # inplace sorting
         unsorted.sort()

         print(unsorted)

[-1.25954441 -0.00846992  0.06725739  0.23700386  0.32812902  0.62523893
  0.79000774  1.01884829  1.43855797  2.59885886]
```

Finding Unique elements:

```
In [43]: array = np.array([1,2,1,4,2,1,4,2])

         print(np.unique(array))

[1 2 4]
```

Set Operations with np.array data type:

```
In [ ]: s1 = np.array(['desk','chair','bulb'])
         s2 = np.array(['lamp','bulb','chair'])
         print(s1, s2)
```

```
In [ ]: print( np.intersect1d(s1, s2) )
```

```
In [ ]: print( np.union1d(s1, s2) )
```

```
In [ ]: print( np.setdiff1d(s1, s2) )# elements in s1 that are not in s2
```

```
In [ ]: print( np.in1d(s1, s2) )#which element of s1 is also in s2
```

Broadcasting:

Introduction to broadcasting.

For more details, please see:

<https://docs.scipy.org/doc/numpy-1.10.1/user/basics.broadcasting.html> (<https://docs.scipy.org/doc/numpy-1.10.1/user/basics.broadcasting.html>)

```
In [80]: import numpy as np

start = np.zeros((4,3))
print(start)

[[ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
 [ 0.  0.  0.]
```

```
In [ ]: # create a rank 1 ndarray with 3 values
add_rows = np.array([1, 0, 2])
print(add_rows)
```

```
In [ ]: y = start + add_rows # add to each row of 'start' using broadcasting
print(y)
```

```
In [83]: # create an ndarray which is 4 x 1 to broadcast across columns
add_cols = np.array([[0,1,2,3]])
add_cols = add_cols.T

print(add_cols)

[[0]
 [1]
 [2]
 [3]]
```

```
In [84]: # add to each column of 'start' using broadcasting
y = start + add_cols
print(y)

[[ 0.  0.  0.]
 [ 1.  1.  1.]
 [ 2.  2.  2.]
 [ 3.  3.  3.]
```

```
In [85]: # this will just broadcast in both dimensions
add_scalar = np.array([1])
print(start+add_scalar)

[[ 1.  1.  1.]
 [ 1.  1.  1.]
 [ 1.  1.  1.]
 [ 1.  1.  1.]
```

Example from the slides:

```
In [86]: # create our 3x4 matrix
arrA = np.array([[1,2,3,4],[5,6,7,8],[9,10,11,12]])
print(arrA)

[[ 1  2  3  4]
 [ 5  6  7  8]
 [ 9 10 11 12]]
```

```
In [87]: # create our 4x1 array
arrB = [0,1,0,2]
print(arrB)

[0, 1, 0, 2]
```

```
In [88]: # add the two together using broadcasting
print(arrA + arrB)

[[ 1  3  3  6]
 [ 5  7  7 10]
 [ 9 11 11 14]]
```

Speedtest: ndarrays vs lists

First setup paramaters for the speed test. We'll be testing time to sum elements in an ndarray versus a list.

```
In [ ]: from numpy import arange
from timeit import Timer

size    = 1000000
timeits = 1000
```

```
In [ ]: # create the ndarray with values 0,1,2...,size-1
nd_array = arange(size)
print( type(nd_array) )
```

```
In [ ]: # timer expects the operation as a parameter,  
# here we pass nd_array.sum()  
timer_numpy = Timer("nd_array.sum()", "from __main__ import nd_array")  
  
print("Time taken by numpy ndarray: %f seconds" %  
      (timer_numpy.timeit(timeits)/timeits))
```

```
In [ ]: # create the list with values 0,1,2...,size-1  
a_list = list(range(size))  
print (type(a_list) )
```

```
In [ ]: # timer expects the operation as a parameter, here we pass sum(a_list)  
timer_list = Timer("sum(a_list)", "from __main__ import a_list")  
  
print("Time taken by list: %f seconds" %  
      (timer_list.timeit(timeits)/timeits))
```

Read or Write to Disk:

Binary Format:

```
In [ ]: x = np.array([ 23.23, 24.24] )
```

```
In [ ]: np.save('an_array', x)
```

```
In [ ]: np.load('an_array.npy')
```

Text Format:

```
In [ ]: np.savetxt('array.txt', X=x, delimiter=',')
```

```
In [ ]: !cat array.txt
```

```
In [ ]: np.loadtxt('array.txt', delimiter=',')
```

Additional Common ndarray Operations

Dot Product on Matrices and Inner Product on Vectors:

```
In [44]: # determine the dot product of two matrices
x2d = np.array([[1,1],[1,1]])
y2d = np.array([[2,2],[2,2]])
```

```
print(x2d.dot(y2d))
print()
print(np.dot(x2d, y2d))
```

```
[[4 4]
 [4 4]]
```

```
[[4 4]
 [4 4]]
```

```
In [45]: # determine the inner product of two vectors
a1d = np.array([9 , 9 ])
b1d = np.array([10, 10])
```

```
print(a1d.dot(b1d))
print()
print(np.dot(a1d, b1d))
```

```
180
```

```
180
```

```
In [ ]: # dot produce on an array and vector
print(x2d.dot(a1d))
print()
print(np.dot(x2d, a1d))
```

Sum:

```
In [78]: # sum elements in the array
ex1 = np.array([[11,12],[21,22]])

print(np.sum(ex1))          # add all members
print(ex1.sum())
```

```
66
```

```
66
```

```
In [ ]: print(np.sum(ex1, axis=0)) # columnwise sum
```

```
In [ ]: print(np.sum(ex1, axis=1)) # rowwise sum
```

Element-wise Functions:

For example, let's compare two arrays values to get the maximum of each.

```
In [46]: # random array
x = np.random.randn(8)
x
```

```
Out[46]: array([-0.84329306,  0.53577379, -0.91005308,  0.19388413,  1.1713282 ,
                1.26538085,  2.90171378,  0.53955055])
```

```
In [47]: # another random array
y = np.random.randn(8)
y
```

```
Out[47]: array([-0.36783814, -0.22151073,  1.76139501,  1.26496062,  0.12611553,
                1.15053687,  0.32329357, -1.64598334])
```

```
In [48]: # returns element wise maximum between two arrays

np.maximum(x, y)
```

```
Out[48]: array([-0.36783814,  0.53577379,  1.76139501,  1.26496062,  1.1713282 ,
                1.26538085,  2.90171378,  0.53955055])
```

Reshaping array:

```
In [49]: # grab values from 0 through 19 in an array
arr = np.arange(20)
print(arr)
```

```
[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19]
```

```
In [50]: # reshape to be a 4 x 5 matrix
arr.reshape(4,5)
```

```
Out[50]: array([[ 0,  1,  2,  3,  4],
                [ 5,  6,  7,  8,  9],
                [10, 11, 12, 13, 14],
                [15, 16, 17, 18, 19]])
```

Transpose:

```
In [51]: # transpose
ex1 = np.array([[11,12],[21,22]])

ex1.T
```

```
Out[51]: array([[11, 21],
               [12, 22]])
```

Indexing using where():

```
In [52]: x_1 = np.array([1,2,3,4,5])

y_1 = np.array([11,22,33,44,55])

filter = np.array([True, False, True, False, True])
```

```
In [53]: out = np.where(filter, x_1, y_1)
print(out)

[ 1 22  3 44  5]
```

```
In [54]: mat = np.random.rand(5,5)
mat
```

```
Out[54]: array([[ 0.67745669,  0.23937408,  0.19180203,  0.04914193,  0.59530901],
               [ 0.81153492,  0.97359325,  0.93238738,  0.78785179,  0.03321135],
               [ 0.35428268,  0.13906719,  0.93602509,  0.93450562,  0.47595539],
               [ 0.82772688,  0.94949964,  0.67273987,  0.34979533,  0.03106072],
               [ 0.83550762,  0.99417402,  0.52919541,  0.30741014,  0.99187241]])
```

```
In [55]: np.where( mat > 0.5, 1000, -1) #ternary operator- if true, 1000 if false, -1.
      Same for previous example.
```

```
Out[55]: array([[1000,  -1,  -1,  -1, 1000],
               [1000, 1000, 1000, 1000,  -1],
               [ -1,  -1, 1000, 1000,  -1],
               [1000, 1000, 1000,  -1,  -1],
               [1000, 1000, 1000,  -1, 1000]])
```

"any" or "all" conditionals:

```
In [65]: arr_bools = np.array([ True, False, True, True, False ])
```

```
In [67]: arr_bools.any() #any of them true in the array
```

```
Out[67]: False
```



```
In [64]: arr_bools.all() #all of them true in the array
```

```
Out[64]: True
```

Random Number Generation:

```
In [68]: Y = np.random.normal(size = (1,5))[0]  
print(Y)
```

```
[-0.06997155  0.59965832  0.06521293  0.19011809  0.16063399]
```

```
In [69]: Z = np.random.randint(low=2,high=50,size=4)  
print(Z)
```

```
[34  2 29 40]
```

```
In [70]: np.random.permutation(Z) #return a new ordering of elements in Z
```

```
Out[70]: array([40, 29, 34,  2])
```

```
In [71]: np.random.uniform(size=4) #uniform distribution
```

```
Out[71]: array([ 0.71689868,  0.06642476,  0.47574044,  0.72704265])
```

```
In [72]: np.random.normal(size=4) #normal distribution
```

```
Out[72]: array([-0.3189061 ,  2.55075888,  0.10701857,  0.99111263])
```

Merging data sets:

```
In [73]: K = np.random.randint(low=2,high=50,size=(2,2))  
print(K)
```

```
print()
```

```
M = np.random.randint(low=2,high=50,size=(2,2))  
print(M)
```

```
[[35  3]  
 [36 23]]
```

```
[[26 45]  
 [14 26]]
```

```
In [74]: np.vstack((K,M))
```

```
Out[74]: array([[35,  3],  
               [36, 23],  
               [26, 45],  
               [14, 26]])
```

```
In [75]: np.hstack((K,M))
```

```
Out[75]: array([[35,  3, 26, 45],  
               [36, 23, 14, 26]])
```

```
In [76]: np.concatenate([K, M], axis = 0)
```

```
Out[76]: array([[35,  3],  
               [36, 23],  
               [26, 45],  
               [14, 26]])
```

```
In [77]: np.concatenate([K, M.T], axis = 1)
```

```
Out[77]: array([[35,  3, 26, 14],  
               [36, 23, 45, 26]])
```

```
In [ ]:
```